

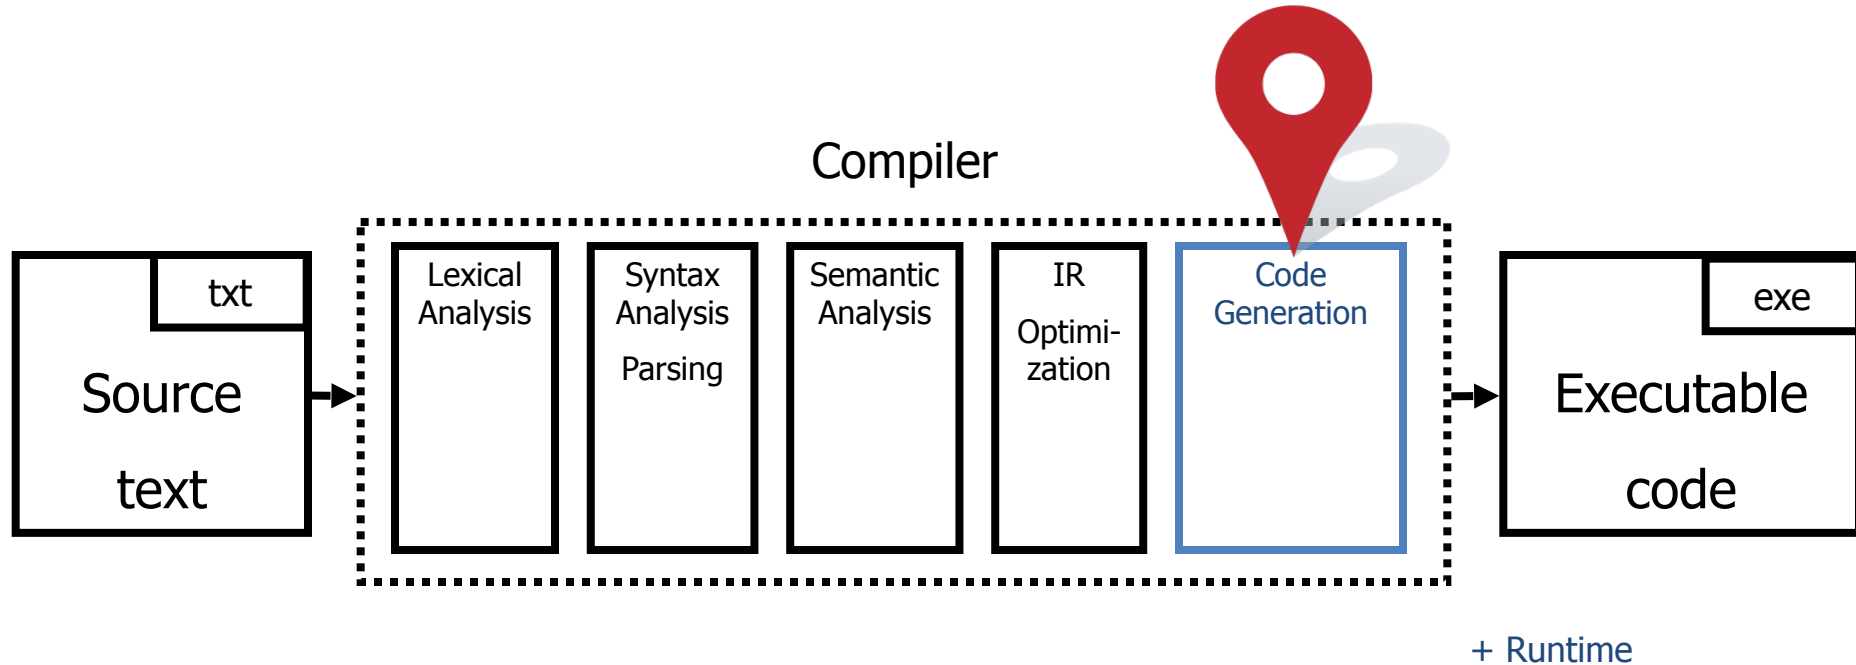
THEORY OF COMPILATION

LECTURE 09



ACTIVATION
RECORDS

You are here



Supporting Procedures

```
n = f(a[i]);
```



```
t1 := i * 4  
t2 := a + t1  
t3 := *t2
```



*code that evaluates
the parameter a[i]*

```
param t3
```

```
t4 := call f, 1
```

of parameters

```
n := t4
```

- How is that done?
- What do we need from the compiler?

Supporting Procedures

- Extending our computing environment
 - ▶ (at least) enough memory for local variables
- Passing information into the new environment
 - ▶ call parameters
- Transfer of control to/from procedure
- Handling return values

Design Decisions

- Scoping rules
 - ▶ Static scoping vs. dynamic scoping
- Memory layout
 - ▶ Allocating space for local variables
- Caller/callee conventions
 - ▶ Saving and restoring register values

Static (Lexical) Scoping

```
main ( )  
{  
  int a = 0;  
  int b = 0;  
  {  
    int b = 1;  
    {  
      int a = 2;  
      printf("%d %d\n", a, b);  
    }  
    {  
      int b = 3;  
      printf("%d %d\n", a, b);  
    }  
    printf ("%d %d\n", a, b);  
  }  
  printf ("%d %d\n", a, b);  
}
```

The diagram illustrates static (lexical) scoping with nested scopes. The scopes are labeled B₀, B₁, B₂, and B₃. B₀ is the outermost scope (main). B₁ is a block-level scope containing a loop body. B₂ is a block-level scope nested within B₁. B₃ is a block-level scope nested within B₁. Arrows indicate that the 'a' in B₂ refers to the 'a' in B₀, and the 'b' in B₃ refers to the 'b' in B₁.


A name refers to its
(closest) enclosing scope

known at compile time

Declaration	Scopes
a = 0	B ₀ , B ₁ , B ₃
b = 0	B ₀
b = 1	B ₁ , B ₂
a = 2	B ₂
b = 3	B ₃

Dynamic Scoping

- Every function call creates new definitions for variables in its scope
 - definitions are maintained in a **global stack**
- When entering scope
 - push local variables on stack
- When exiting scope
 - pop local variables
- **Evaluating the identifier in any context binds to the current top of stack**
 - ▶ determined **at runtime**



Example

```
int x = 42;  
  
int f() { return x; }  
int g() { int x = 1; return f(); }  
int main() { print g(); print x; }
```

- What values are output by main?
 - ▶ static scoping?
 - ▶ dynamic scoping?

Why do we care?

- **We need to generate code to access variables**
- Static scoping
 - identifier binding is known at **compile time**
 - address of the variable is known at **compile time**
 - assigning addresses to variables is part of code generation
 - no runtime errors of “access to undefined variable”
 - can check types of variables

Variable Addressing for Static Scoping (first attempt)

```
int x = 42;
```

```
int f() { return x; }
```

```
int g() { int x = 1;  
        return f(); }
```

```
int main() { print g(); }
```

identifier	address
x (global)	0x42
x (inside g)	0x73

Variable Addressing for Static Scoping (first attempt)

```
int a[11];

void quicksort(int m, int n)
{
    int i;
    if (n > m) {
        i = partition(m, n);
        quicksort(m, i-1);
        quicksort(i+1, n);
    }
    :

int main() {
    ...
    quicksort(1, 9);
}
```

identifier	address
a (global)	0x42
i (inside quicksort)	...

what is the address of the
variable **i** in the procedure
quicksort?

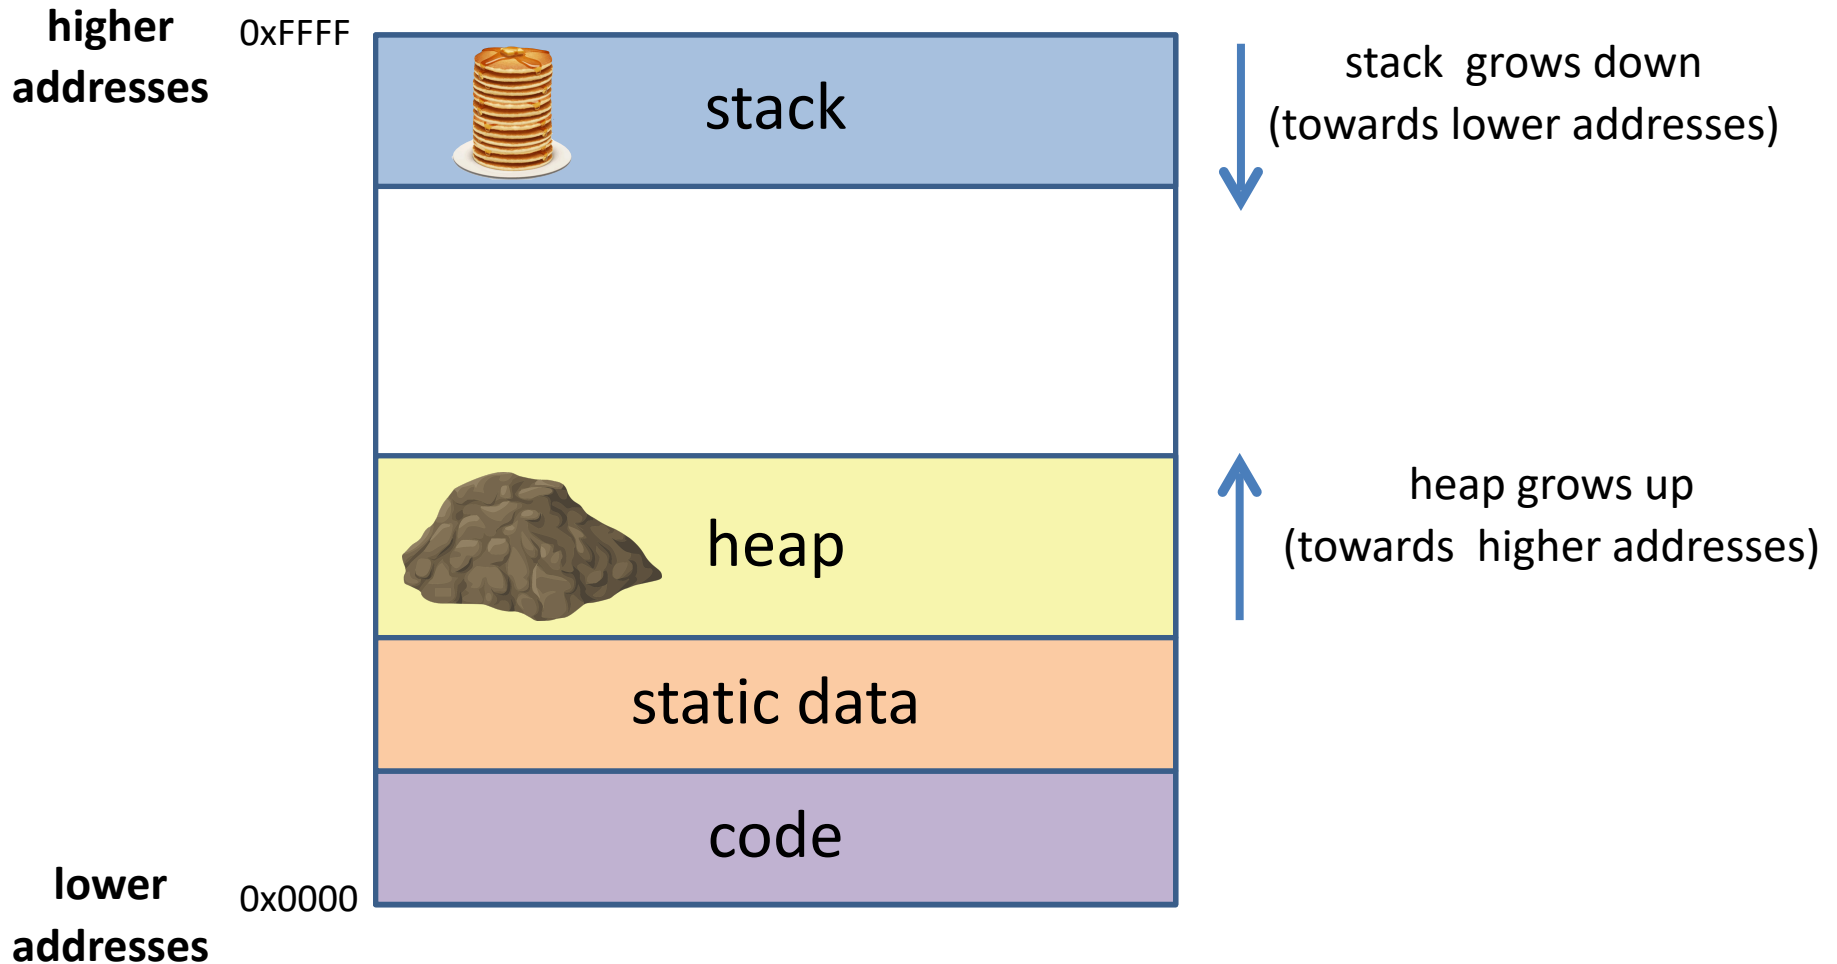
(or arguments **m**, **n**,
for that matter?)

Activation Record (frame)

- Separate space for each procedure invocation
 - ▶ **Managed at runtime**
 - ▶ But **code** for managing it generated **by the compiler**
- Desired properties
 - ▶ Efficient allocation and deallocation
 - procedure calls are frequent
 - ▶ Flexible size
 - different procedures may require different memory sizes




Memory Layout



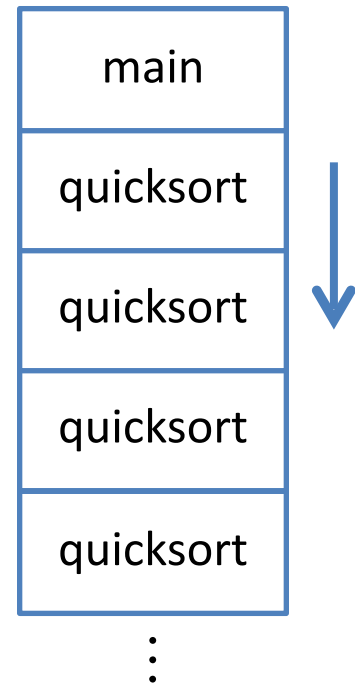
Runtime Stack

- FP – frame pointer
 - beginning of current frame + some offset
 - Sometimes called BP (base pointer)
- SP – stack pointer
 - end of current frame



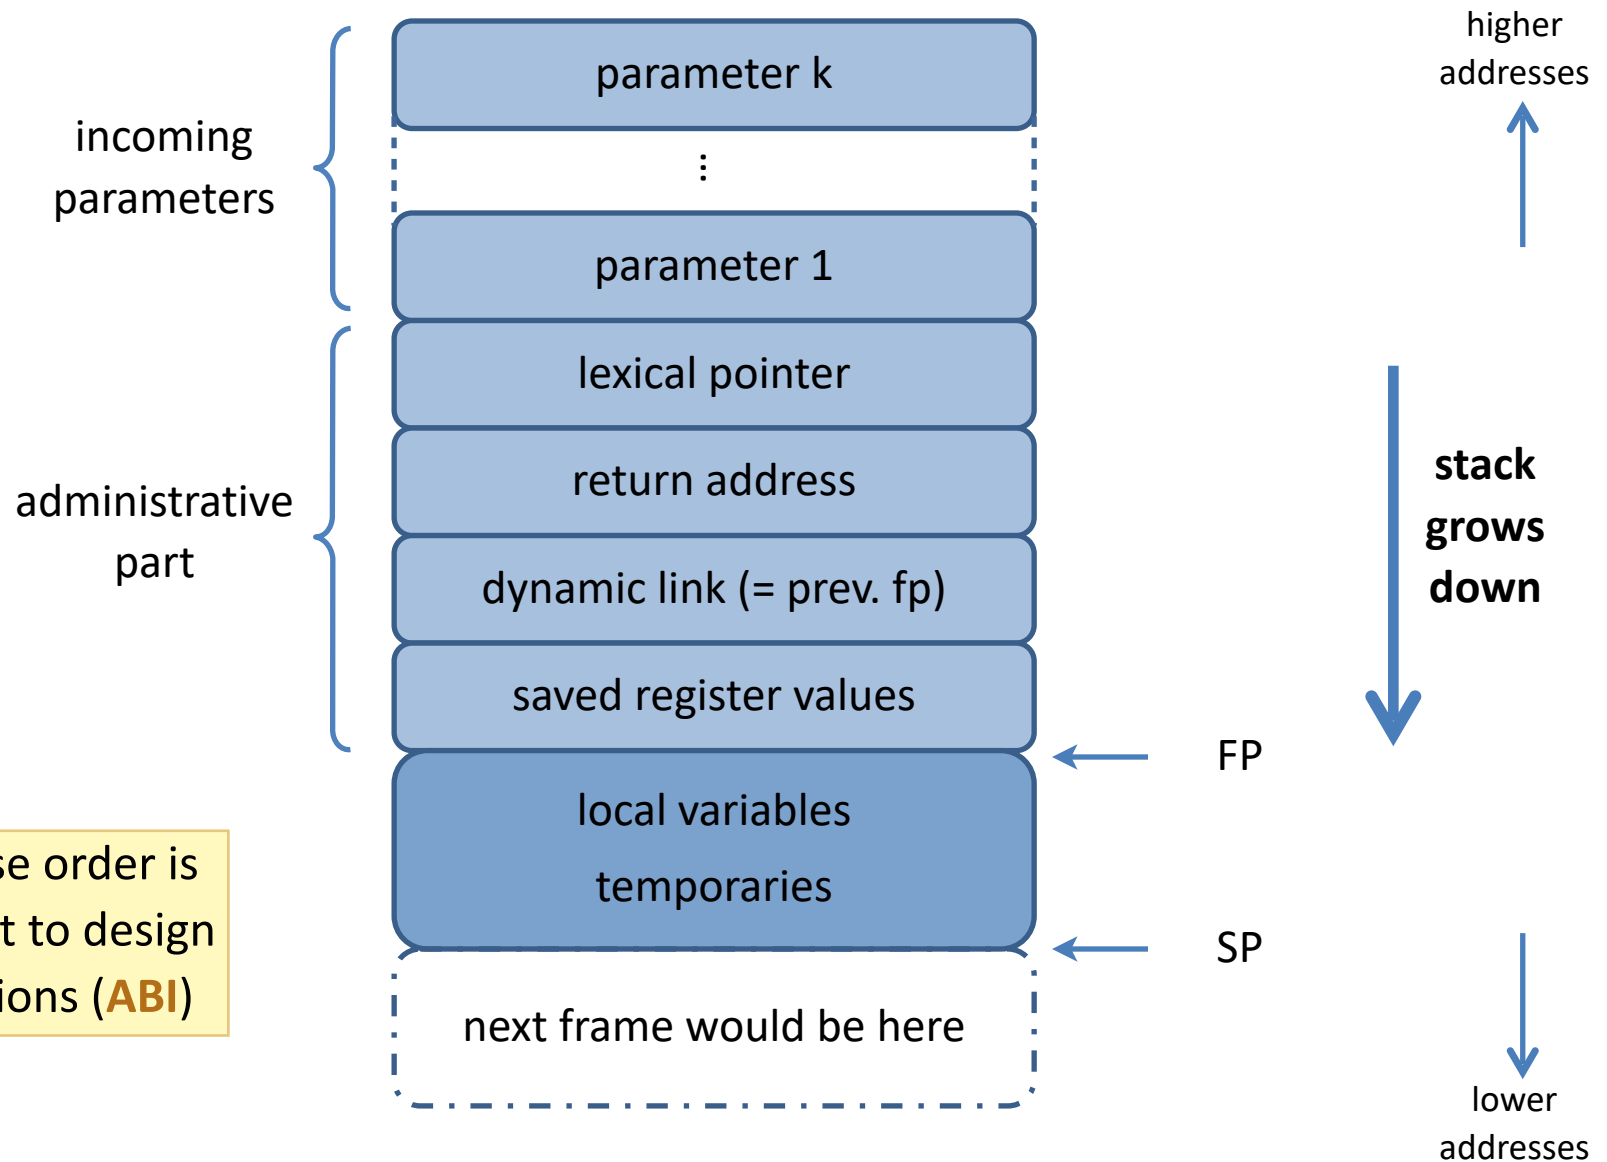
Runtime Stack

- Stack of activation records
 - ▶ Call = push new activation record
 - ▶ Return = pop activation record
- Only one “active” record at a time*
 - ▶ top of stack
- This allows to handle recursion



* per thread

Activation Record (frame)



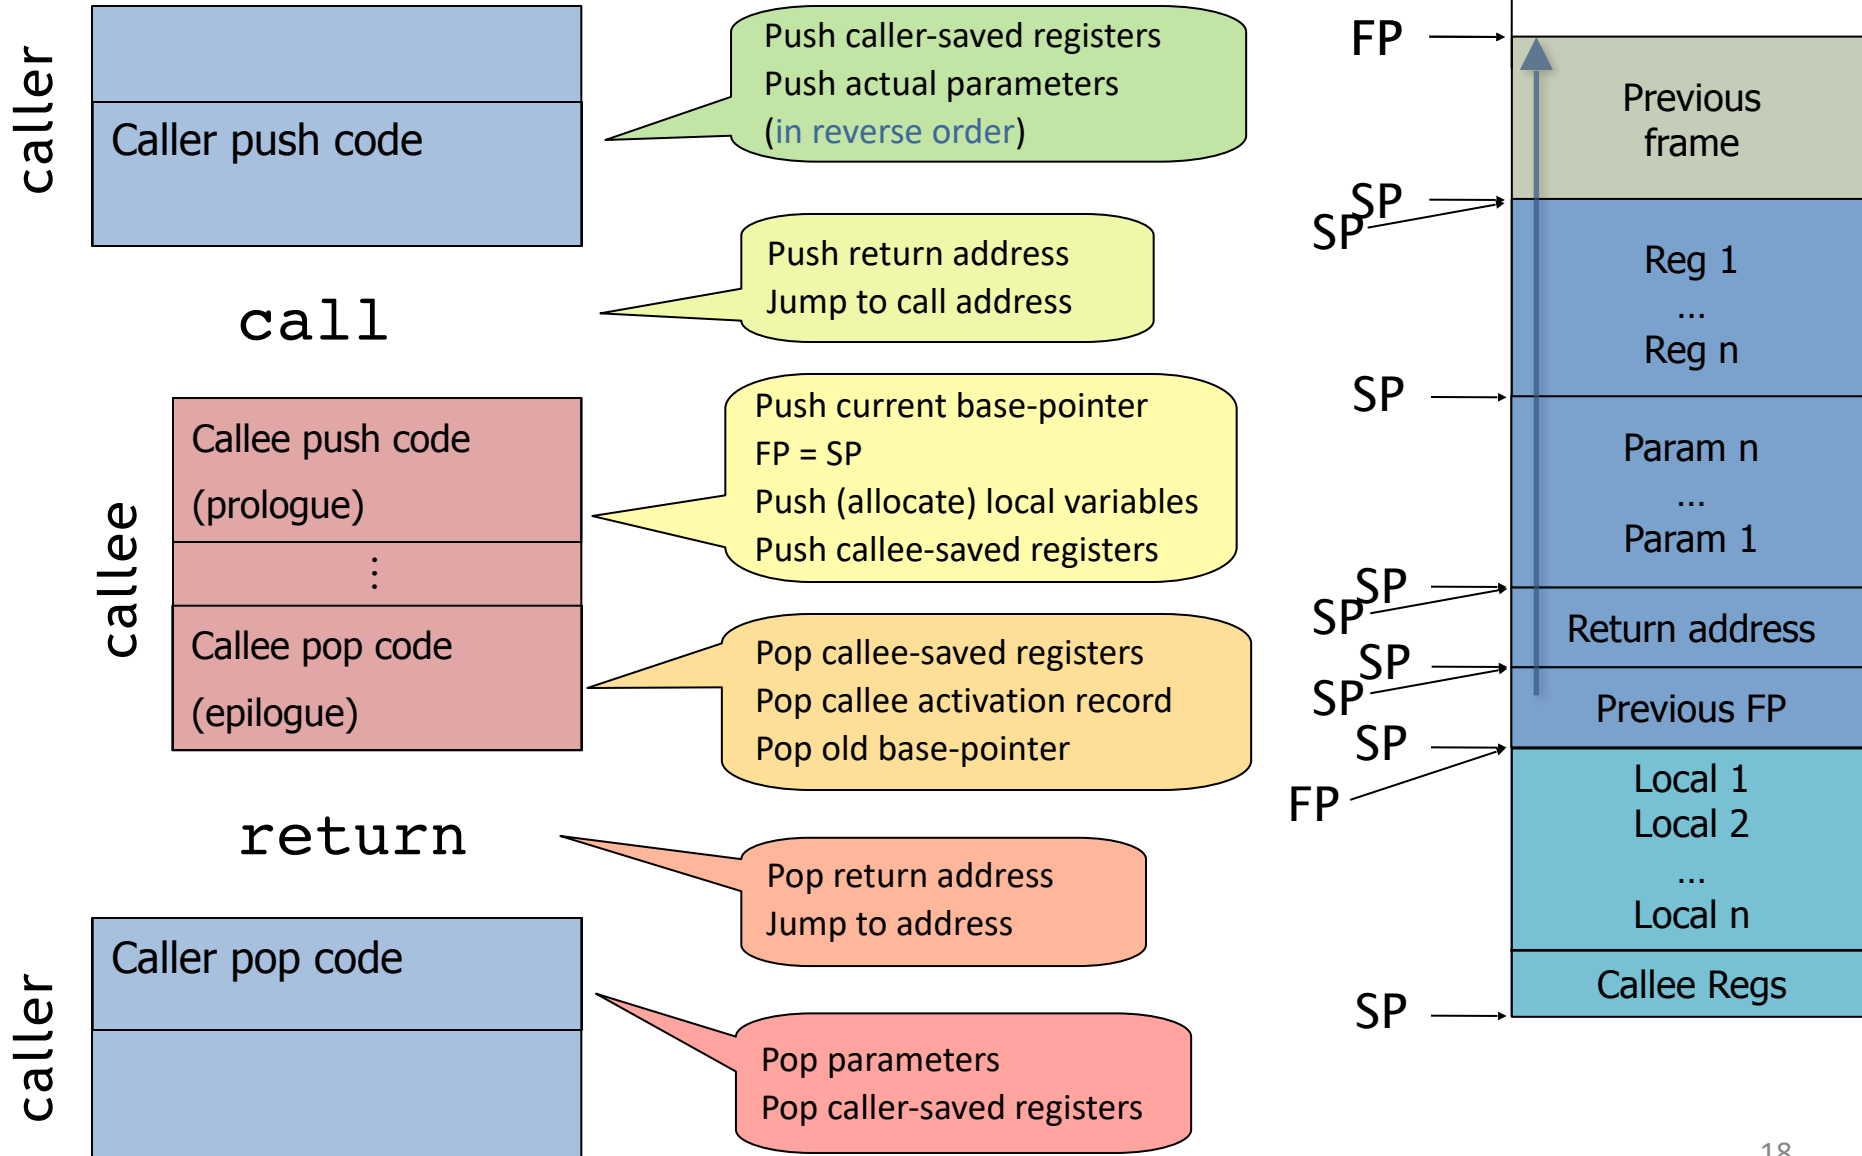
Call Sequences

- The processor **does not** save the content of registers on “call” instruction
 - ▶ (Why is this a problem?)
- So who will?

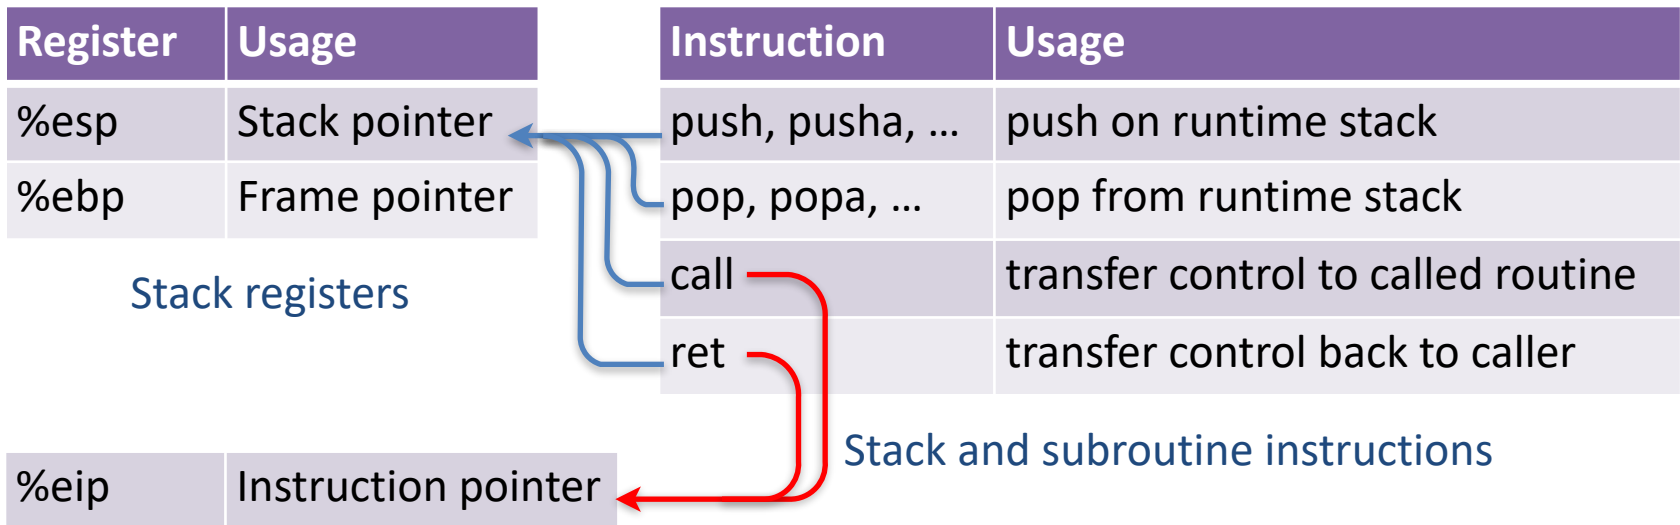
There are several options:

 - ▶ **Caller** saves registers before call, restores after call
 - ▶ **Callee** saves when called, restores before return
 - ▶ But can also have **both** save/restore some registers

Call Sequences



x86 Runtime Stack



foo(42, 21)

Call Sequences

x86 (32 bit), cdecl calling convention

caller

```
push %ecx
push $21
push $42
call _foo
```

call

```
push %ebp
mov %esp, %ebp
sub $8, %esp
push %ebx
:
```

callee

```
pop %ebx
mov %ebp, %esp
pop %ebp
ret
```

return

caller

```
add $8, %esp
pop %ecx
```

Push caller-saved registers
Push actual parameters
(in reverse order)

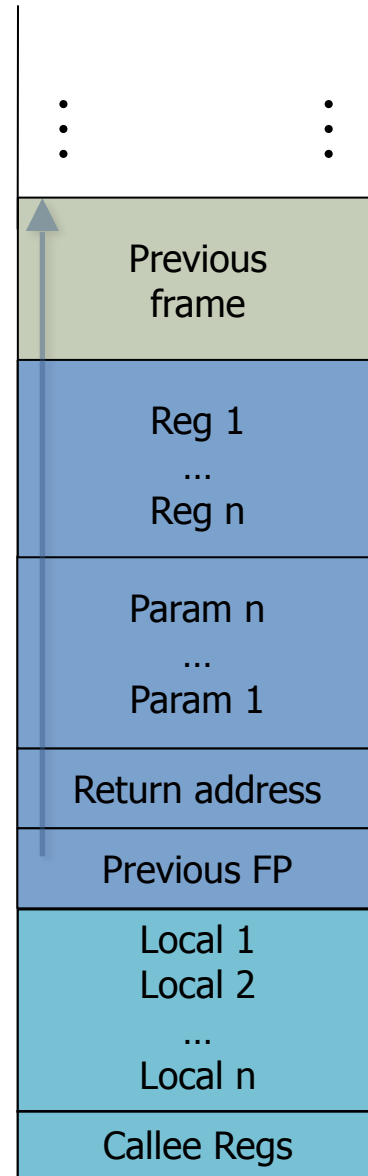
Push return address
Jump to call address

Push current base-pointer
FP = SP
Push (allocate) local variables
Push callee-saved registers

Pop callee-saved registers
Pop callee activation record
Pop old base-pointer

Pop return address
Jump to address

Pop parameters
Pop caller-saved registers



“To Callee-save or to Caller-save?”

- That is indeed the question
 - ▶ **Callee-saved** registers need only be saved when callee modifies their value
 - ▶ **Caller-saved** registers need only be saved if the caller needs their value after the call returns
- Some conventions and heuristics are followed
 - ▶ x86 cdecl: `%eax, %ecx, %edx` are **caller-saved**; the rest are **callee-saved**.

These **calling conventions** are part of the architecture's **ABI**
(Application Binary Interface)

“To Callee-save or to Caller-save?”

x86 cdecl: `%eax, %ecx, %edx` are **caller-saved**;
the rest are **callee-saved**.

bar:

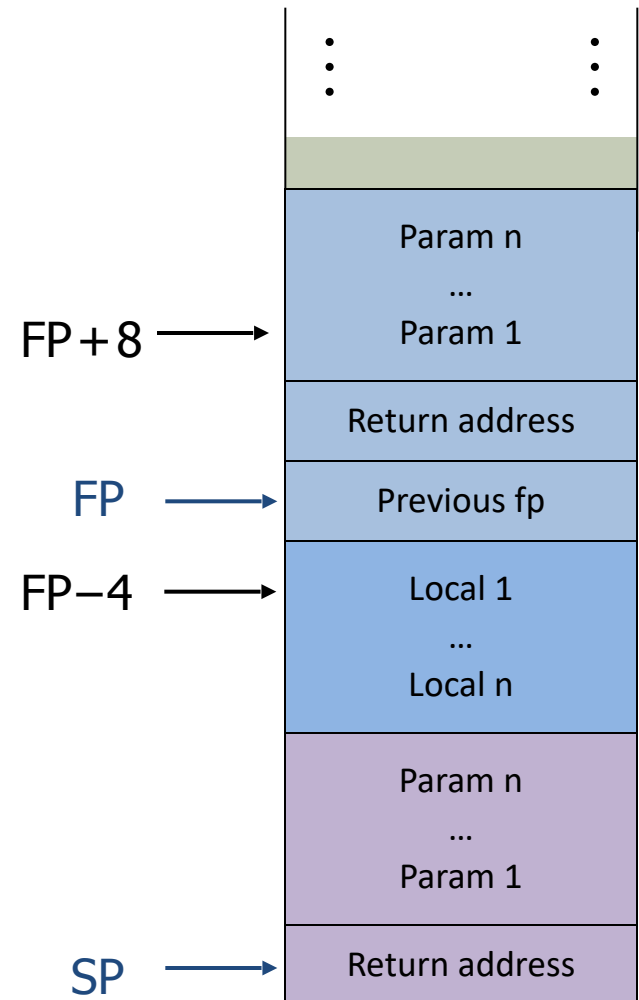
```
push %ebp
mov %esp, %ebp
push %ebx
mov %ebx, 6
mov (@garray, %ebx, 4), %eax
mov %ebx, 7
mov (@garray, %ebx, 4), %ecx
add %ecx, %eax
pop %ebx
pop %ebp
ret
```

main:

```
printf("%d", bar());
:
mov %edx, (@str) ; "%d"
push %edx
call bar
pop %edx
push %eax
push %edx
call printf
:
```

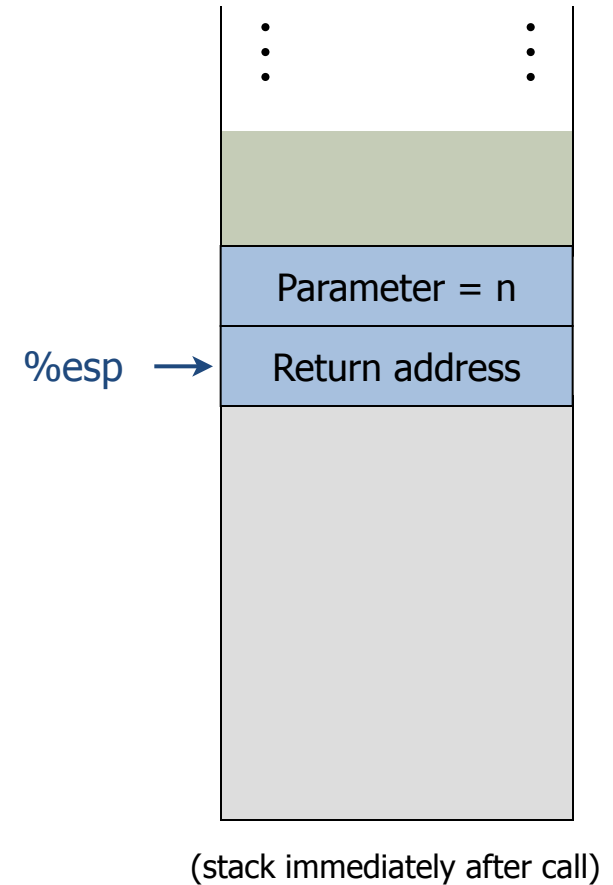

Accessing Stack Variables

- Use offset from FP (%ebp)
- Remember – stack grows downwards
- Above FP = parameters
- Below FP = locals
- Examples
 - $FP + 4$ = return address
 - $FP + 8$ = first parameter
 - $FP - 4$ = first local



Factorial – `fact (int n)`

```
int fact(int n) {  
    if (n <= 1) {  
        return 1;  
    }  
    else {  
        return fact(n - 1) * n;  
    }  
}
```



Factorial — `fact(int n)`

`fact:`

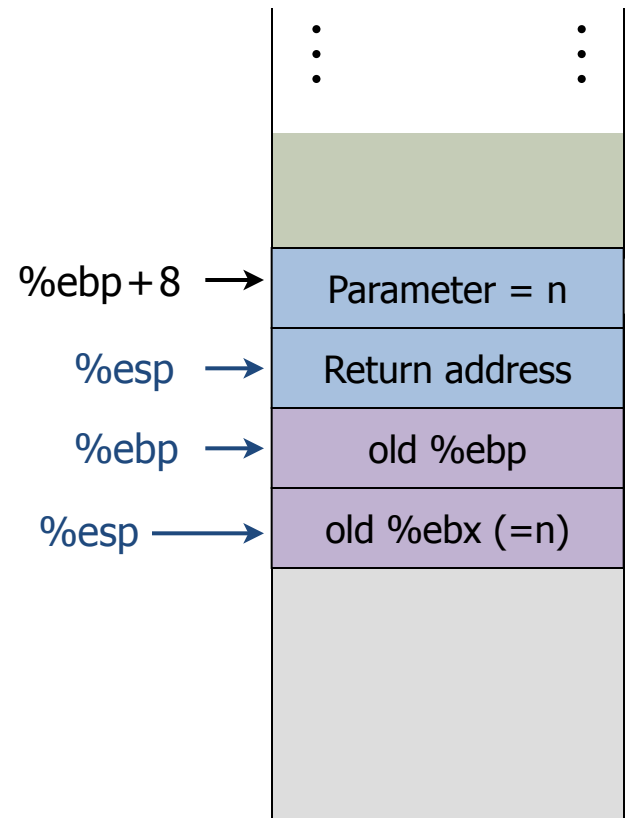
```
pushl %ebp          # save ebp
movl %esp,%ebp      # ebp = esp
pushl %ebx          # save ebx
movl 8(%ebp),%ebx    # ebx = n
cmpl $1,%ebx        # n <= 1 ?
jle .lresult        # then done
leal -1(%ebx),%eax   # eax = n-1

pushl %eax          #
call fact           # eax = fact(n-1)
addl $4,%esp        #

imull %ebx,%eax      # eax = eax * n
jmp .lreturn        #

.lresult:
movl $1,%eax        # return 1

.lreturn:
movl -4(%ebp),%ebx   # restore ebx
movl %ebp,%esp       # restore esp
popl %ebp           # restore ebp
ret
```



(stack after prologue)

Factorial — `fact(int n)`

`fact:`

```

pushl %ebp          # save ebp
movl %esp,%ebp      # ebp = esp
pushl %ebx          # save ebx
movl 8(%ebp),%ebx    # ebx = n
cmpl $1,%ebx        # n <= 1 ?
jle .lresult        # then done
leal -1(%ebx),%eax   # eax = n-1

```

```

pushl %eax          #
call fact           # eax = fact(n-1)

```

```

( addl $4,%esp )    #
imull %ebx,%eax     # eax = n * fact(n-1)
jmp .lreturn        #

```

`lresult:`

```

movl $1,%eax        # return 1

```

`lreturn:`

```

movl -4(%ebp),%ebx  # restore ebx
movl %ebp,%esp      # restore esp
popl %ebp           # restore ebp
ret

```

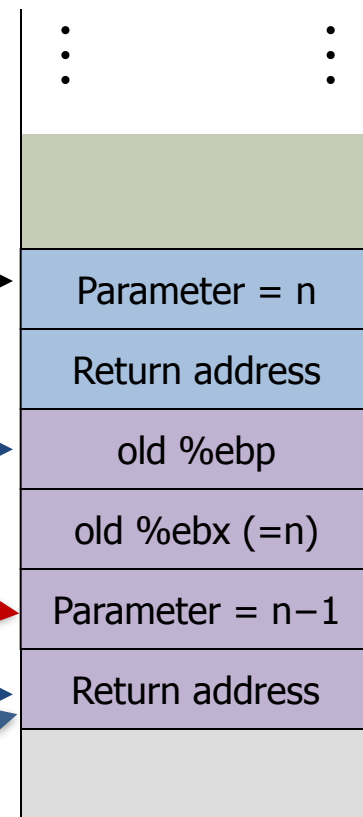
Caller pops
parameters

Callee pops
return address

`%ebp+8` →

`%ebp` →

`%esp` →



(stack immediately after second call)

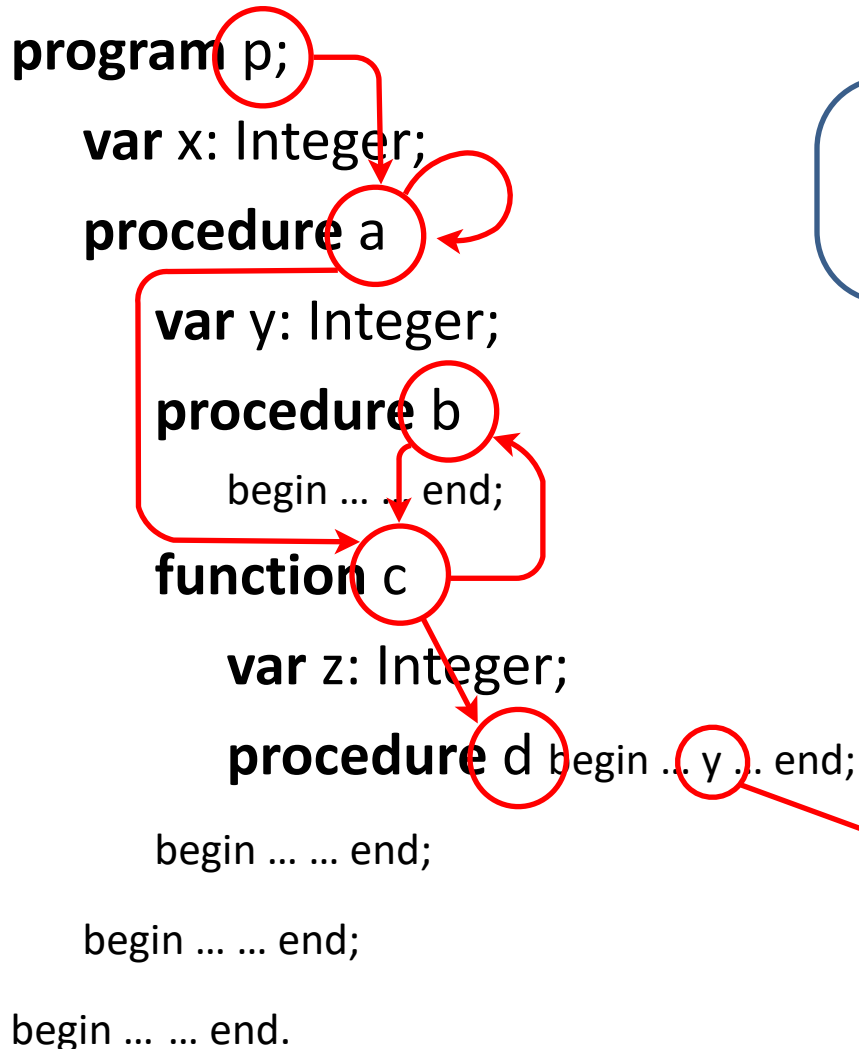
Nested Procedures

Exist *e.g.* in Pascal, JavaScript

- Idea: A routine can have sub-routines
- Any sub-routine can access anything that is defined in its own scope or any surrounding scope
 - “non-local” variables
- A sub-routine is allowed to call itself, its own sub-routines, any of its ancestors, and *also its siblings*
 - (but not “nephews” or “grandchildren”)

Example: Nested Procedures

```
program p;  
  var x: Integer;  
  procedure a  
    var y: Integer;  
    procedure b  
      begin ... end;  
    function c  
      var z: Integer;  
      procedure d begin ... y ... end;  
    begin ... end;  
  begin ... end;  
begin ... end.
```



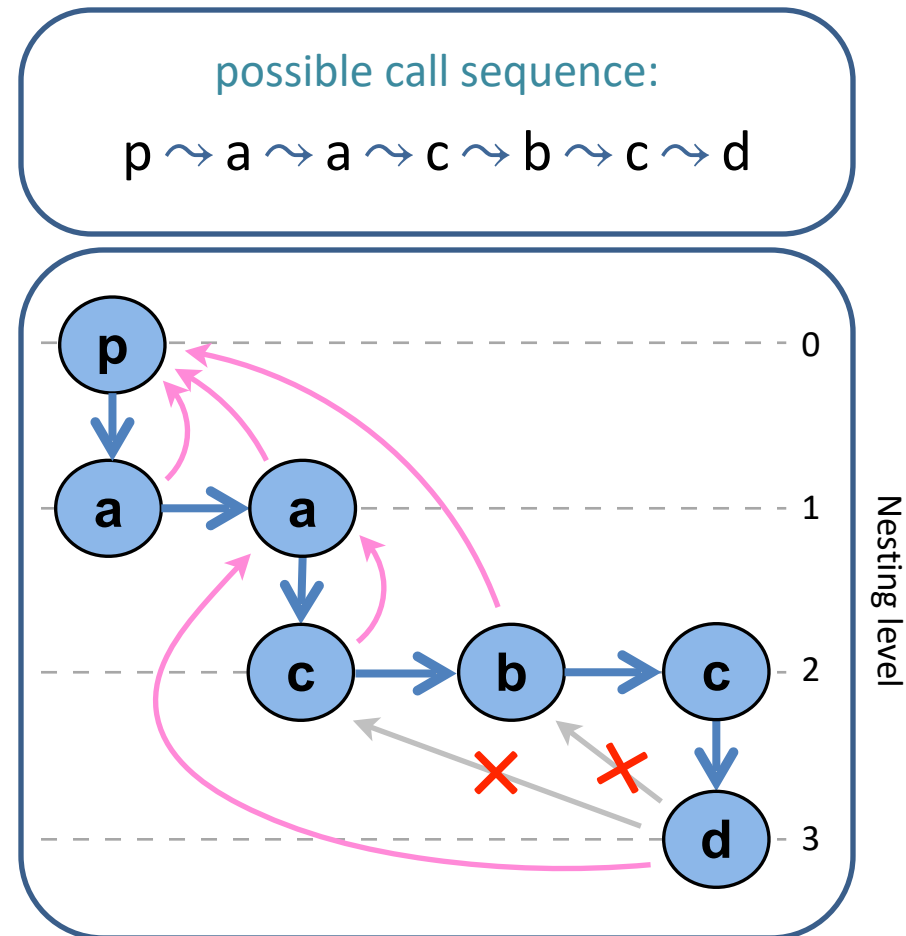
possible call sequence:

$p \rightsquigarrow a \rightsquigarrow a \rightsquigarrow c \rightsquigarrow b \rightsquigarrow c \rightsquigarrow d$

what is the address of
variable “y” in procedure d?

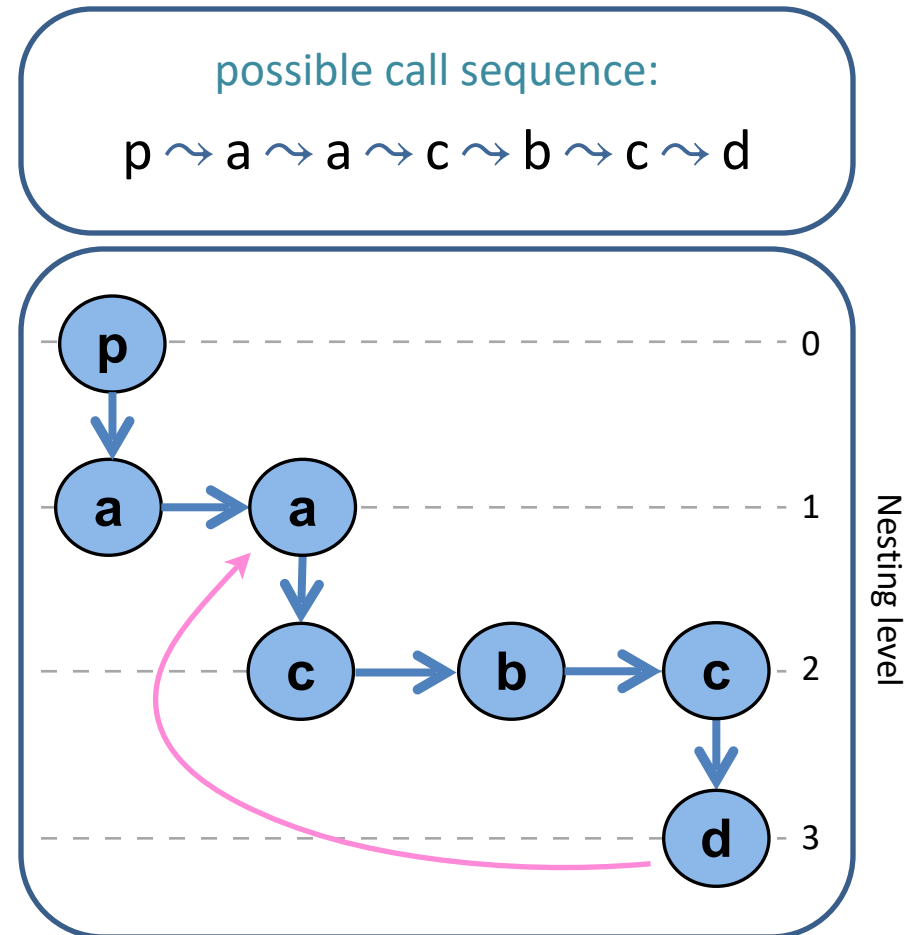
Nested Procedures

- A routine can call a nestling, an ancestor, or a sibling
 - ▶ When “c” uses (non-local) variables from “a”, which instance of “a” is it?
 - ▶ How do you find the right activation record at runtime?



Nested Procedures

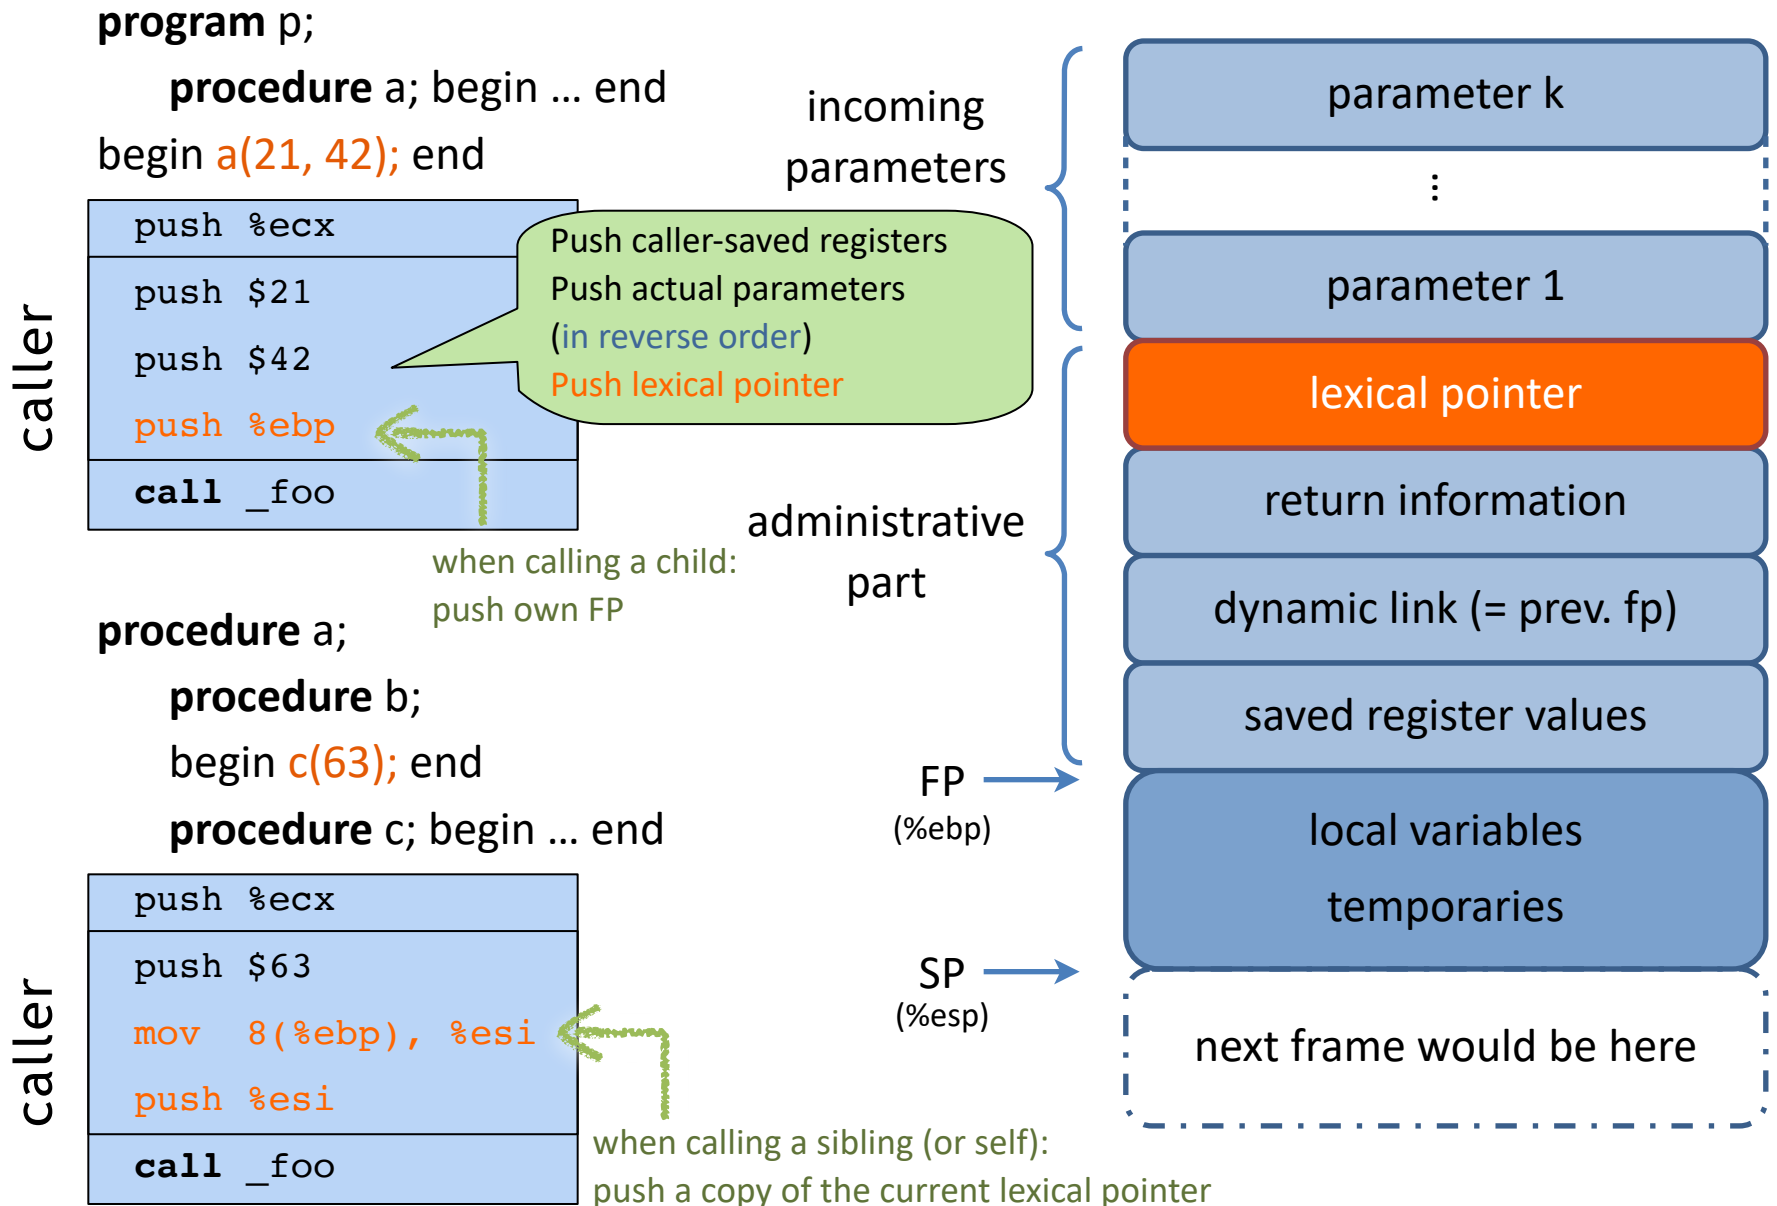
- Goal: find the closest frame of some routine in the stack from a given nesting level
- A routine may occur more than once in a sequence of calls
 - ▶ if a routine at level k uses variables of the same nesting level \Rightarrow it refers to the current activation record
 - ▶ if it uses variables of nesting level $j < k \Rightarrow$ it must be the **last** routine called at level j
- If some routine is last at level j on the stack, then it **must be an ancestor** of the current routine



Nested Procedures

- Goal: find the closest frame of some routine in the stack from a given nesting level
- Solution: **store a *lexical pointer* (a.k.a. *access link*) in the activation record**
 - ▶ Lexical pointer points to the **last** activation record of the nesting level **above** it
 - According to what we just said, this is enough to reach any variable being referenced
- Lexical pointers are created **at runtime**
- Number of links to be traversed is always **known at compile time**

Activation Record (frame)




Lexical Pointers


```
program p;  
  var x: Integer;  
  procedure a  
    var y: Integer;  
    procedure b  
      begin... .. end;  
    function c  
      var z: Integer;  
      procedure d  
        begin... y ... end;  
      begin ... .. end;  
    begin ... .. end;  
  begin ... .. end.
```

possible call sequence:

$p \rightsquigarrow a \rightsquigarrow a \rightsquigarrow c \rightsquigarrow b \rightsquigarrow c \rightsquigarrow d$



Runtime stack



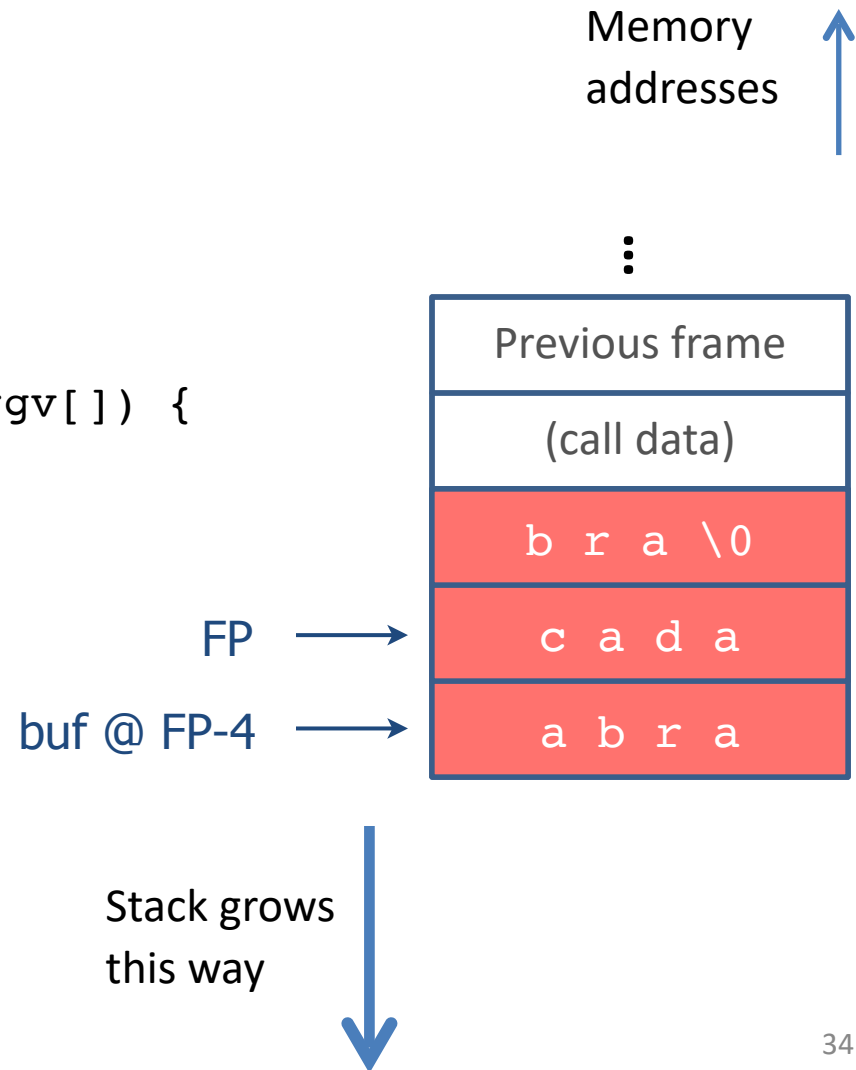
Security Exploit(s)

Buffer Overflow

```
void foo (char *x) {  
    char buf[4];  
    strcpy(buf, x);  
}  
  
int main (int argc, char *argv[]) {  
    foo(argv[1]);  
}
```

```
% ./a.out abracadabra  
Segmentation fault
```

(YMMV)



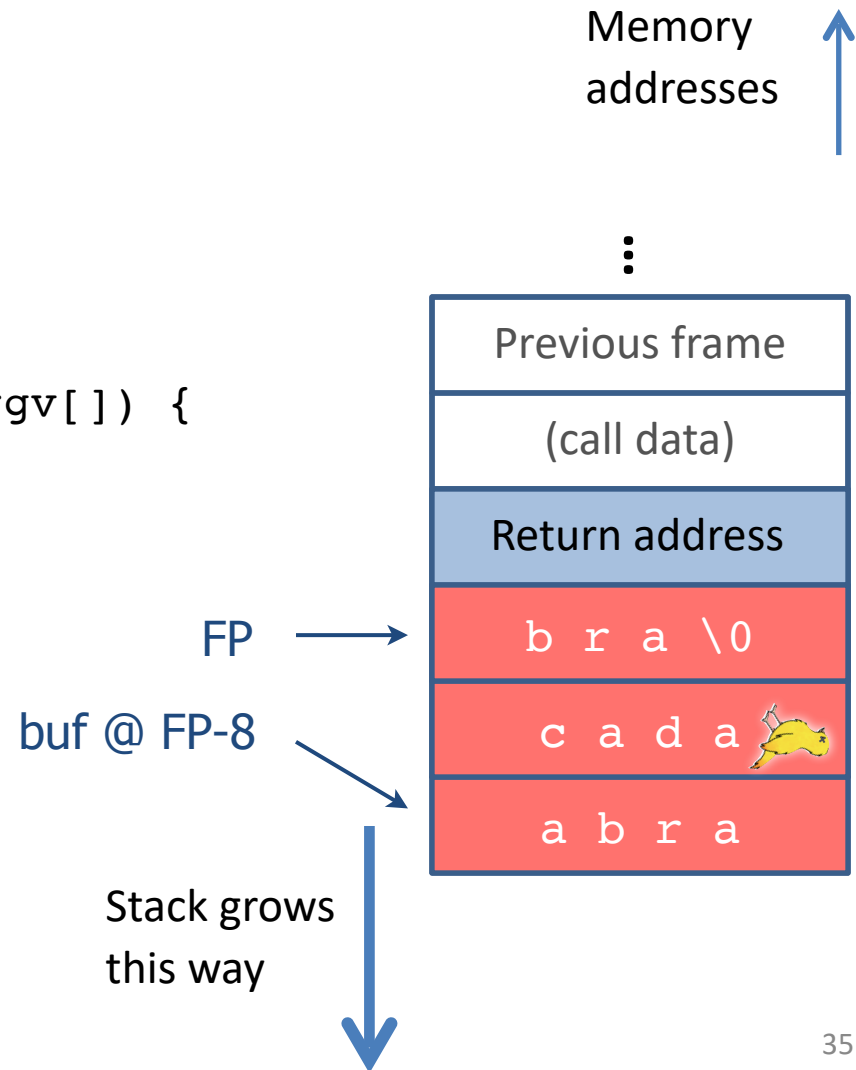
Security Exploit(s)

Buffer Overflow

```
void foo (char *x) {  
    char buf[4];  
    strcpy(buf, x);  
}  
  
int main (int argc, char *argv[]) {  
    foo(argv[1]);  
}
```

```
% ./a.out abracadabra  
<busted>
```

(YMMV)



Buffer overflow

```
int check_authentication(char *password) {
    int auth_flag = 0;
    char pw_buf[16];

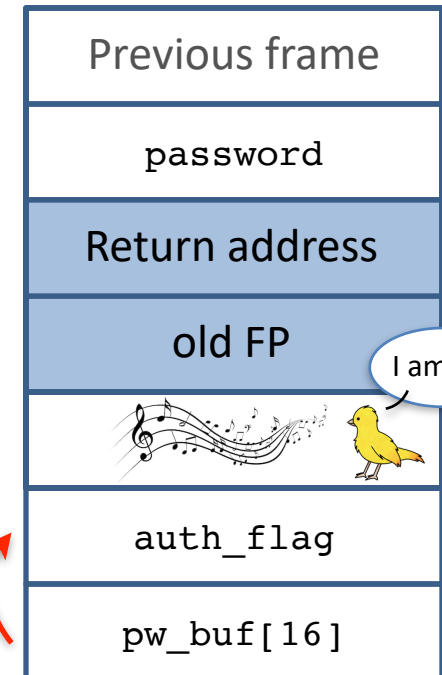
    strcpy(pw_buf, password);
    if (strcmp(pw_buf, "brillig") == 0)
        auth_flag = 1;
    if (strcmp(pw_buf, "outgrabe") == 0)
        auth_flag = 1;
    return auth_flag;
}

int main(int argc, char *argv[]) {
    if (argc < 2) {
        printf("Usage: %s <password>\n", argv[0]);
        exit(0);
    }
    if (check_authentication(argv[1])) {
        printf("\n-----\n");
        printf("        Access Granted.\n");
        printf("-----\n");
    }
    else
        printf("\nAccess Denied.\n");
}
```

Memory
addresses

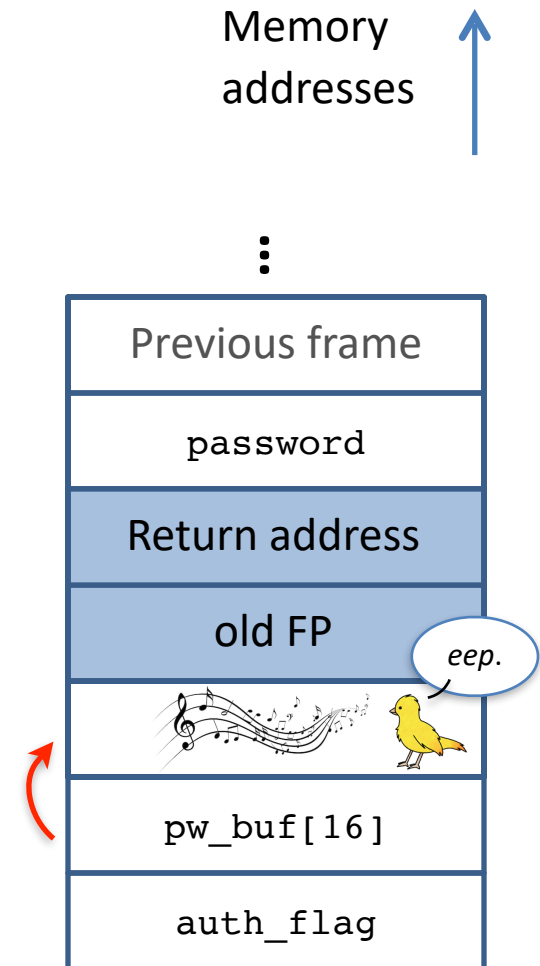


⋮



Buffer overflow

```
int check_authentication(char *password) {  
    char pw_buf[16];  
    int auth_flag = 0;  
  
    strcpy(pw_buf, password);  
    if (strcmp(pw_buf, "brillig") == 0)  
        auth_flag = 1;  
    if (strcmp(pw_buf, "outgrabe") == 0)  
        auth_flag = 1;  
    return auth_flag;  
}  
  
int main(int argc, char *argv[]) {  
    if (argc < 2) {  
        printf("Usage: %s <password>\n", argv[0]);  
        exit(0);  
    }  
    if (check_authentication(argv[1])) {  
        printf("\n-----\n");  
        printf("        Access Granted.\n");  
        printf("-----\n");  
    }  
    else  
        printf("\nAccess Denied.\n");  
}
```



Activation Records: Summary

- ✓ Compile time memory management for procedure data
 - Used to pass parameters, store local variables, and restore registers
- ✓ Works well for data with well-scoped lifetime
 - deallocation when procedure returns
 - no need to call `free()` for stack data

Coming Up

